

# Udacity Self-Driving Car Engineer Nanodegree Highlights

- [Term 1](#)
- [Term 2](#)
- [Term 3](#)
- [References](#)

## Term 1

(Dec 12, 2016 - Apr 4, 2017)

Lesson	Brief	Notes
<b>Introduction</b>		
1	Welcome	<ul style="list-style-type: none"><li>• Curriculum: <a href="https://medium.com/self-driving-cars/term-1-in-depth-on-udacitys-self-driving-car-curriculum-ffcf46af08#.api26st07">https://medium.com/self-driving-cars/term-1-in-depth-on-udacitys-self-driving-car-curriculum-ffcf46af08#.api26st07</a></li><li>• Two approaches to self-driving (both actively pursued):<ul style="list-style-type: none"><li>• robotics: model based sensor fusion and navigation</li><li>• deep learning: learn by mimicking human driving behaviour</li></ul></li><li>• History: 'The Great Robot Race' (below) on how teams adopted different approaches to 2003/2004 <a href="#">DARPA Grand Challenges</a> . Basically, a few focused on hardware as the main challenge while others focused on software. Software won, and by a big margin! The video is long but worth watching. <a href="https://www.youtube.com/watch?v=saV_X9GfIM">https://www.youtube.com/watch?v=saV_X9GfIM</a></li></ul>
2	Finding Lane Lines (project due 22 Dec)	<ul style="list-style-type: none"><li>• Review of Canny Edge Detection and Hough Transforms</li><li>• Lane detection project using OpenCV and Python: <a href="https://www.youtube.com/watch?v=0BG9CxDJ3qk">https://www.youtube.com/watch?v=0BG9CxDJ3qk</a></li></ul>
<b>Deep Learning</b>		
3	Introduction to Deep Learning Module	
4	Regression and Classification	<ul style="list-style-type: none"><li>• Regression by least-squares curve fitting</li><li>• Cross-validation to detect underfitting or overfitting model to data</li></ul>
5	Neural Networks	<ul style="list-style-type: none"><li>• Perceptrons<ul style="list-style-type: none"><li>• compute operations on boolean data (AND, OR, NOT, XOR)</li><li>• Linear and nonlinear separability</li><li>• training: perceptron rule, gradient descent rule</li></ul></li><li>• Back-propagation</li><li>• Types of inductive bias - preference bias and restriction bias</li></ul>
6	Miniflow	<ul style="list-style-type: none"><li>• Hands-on exercise - implementation of a TensorFlow-like system in Python</li></ul>
7	Introduction to Tensorflow	<ul style="list-style-type: none"><li>• Multinomial logistic classifiers: Softmax function, 1-hot encoding, cross-entropy</li><li>• Initialisation and numerical stability of loss function calculations</li><li>• Stochastic Gradient Descent: Momentum and Learning Rate Decay</li></ul>
8	Deep Neural Networks	<ul style="list-style-type: none"><li>• Rectified Linear Units (RELU)</li><li>• Chain rule to setup back-propagation</li><li>• Deep vs wide NNs - advantages</li><li>• Regularisation to prevent overfitting and to increase robustness: L2 regularisation, 'Drop-Out'</li><li>• Implementing deep NNs in TensorFlow</li></ul>
9	Convolutional Neural Networks	<ul style="list-style-type: none"><li>• Weight sharing</li><li>• How CNNs transform input from spatial to semantic information</li><li>• Dimensionality equation (number of neurons per layer)</li><li>• Improving convnets: Pooling, 1x1 convolutions, Inception module</li><li>• Exercise: Implement Lenet-5 deep neural network model (<a href="#">lecun-98.pdf</a>) in TensorFlow</li></ul>

10	Traffic Sign Classifier (Project Due 23 Jan)	<ul style="list-style-type: none"> <li>Traffic sign recognition using multi-scale convolutional networks (<a href="#">sermanet-ijcnn-11.pdf</a>)</li> <li>Developed a network with 197415 parameters, trained on German traffic signs and tested on UK traffic signs</li> <li>Report: <a href="https://github.com/cvillas/CarND/blob/master/P2-TrafficSigns/Traffic_Sign_Classifier.ipynb">https://github.com/cvillas/CarND/blob/master/P2-TrafficSigns/Traffic_Sign_Classifier.ipynb</a></li> </ul>
11	Keras	Re-implemented the traffic sign classifier using both sequential and functional APIs provided by Keras
12	Transfer Learning	<ul style="list-style-type: none"> <li>The idea of re-purposing an existing trained network for a new similar task</li> <li>Feature Extraction and Fine Tuning</li> <li>Repurposed AlexNet (60 million parameters, 650000 neurons) for traffic sign classification</li> <li>Bottleneck features</li> <li>Benchmark VGGNet, GoogLeNet, ResNet on cifar10 and traffic signs datasets</li> <li>The importance of experimenting with network architectures (practise is ahead of theory at the moment)</li> </ul>
13	Behavioral Cloning (Project Due 13 Feb)	<ul style="list-style-type: none"> <li>End-to-end learning (i.e. no modeling or rules) and its advantages (nvidia paper: <a href="#">nvidia_end2end-16.pdf</a>, <a href="#">nvidia_end2end2.pdf</a>)</li> <li>Simulator: <a href="https://techcrunch.com/2017/02/08/udacity-open-sources-its-self-driving-car-simulator-for-anyone-to-use/">https://techcrunch.com/2017/02/08/udacity-open-sources-its-self-driving-car-simulator-for-anyone-to-use/</a></li> <li>Report: <a href="https://github.com/cvillas/CarND/blob/master/P3-BehavioralCloning/submission/report/report.md">https://github.com/cvillas/CarND/blob/master/P3-BehavioralCloning/submission/report/report.md</a></li> <li>My architecture has 546619 parameters to tune. Others have achieved the same using SqueezeNet with 52 parameters: <a href="https://github.com/mez/carnd/blob/master/P3_behavioral_cloning/writeup_report.md">https://github.com/mez/carnd/blob/master/P3_behavioral_cloning/writeup_report.md</a></li> <li>Batch-normalisation in improving learning ability of a network: <a href="#">batch_normalisation.pdf</a></li> <li>Project Video: <a href="https://www.youtube.com/watch?v=WBbD0AUHC_c">https://www.youtube.com/watch?v=WBbD0AUHC_c</a></li> </ul>
<b>Computer Vision</b>		
14	Advanced Lane Finding (Project Due 27 Feb)	<ul style="list-style-type: none"> <li>Reversing lens distortions</li> <li>Perspective transforms</li> <li>Project Video: <a href="https://www.youtube.com/watch?v=zgabVNMf4I8">https://www.youtube.com/watch?v=zgabVNMf4I8</a></li> </ul>
15	Vehicle Detection and Tracking (Project Due 13 Mar)	<ul style="list-style-type: none"> <li>Supervised classification algorithms - naive bayes, support vector machines, decision trees</li> <li>Entropy and information gain</li> <li>Bias-variance dilemma</li> <li>Histogram of Gradients: <a href="#">dalal-cvpr05.pdf</a></li> <li>Creating feature vectors with HoG</li> <li>Multiscale HoG windows</li> <li>Heat maps to combine multiple detections</li> <li>Others have used SSD classifier instead for faster detection (see <a href="https://arxiv.org/abs/1512.02325">https://arxiv.org/abs/1512.02325</a>)</li> <li>Project Video: <a href="https://www.youtube.com/watch?v=D8gm7y9t-7E">https://www.youtube.com/watch?v=D8gm7y9t-7E</a></li> </ul>

## Term 2

(Apr 6, 2017 - Jul 17, 2017)

Lesson	Brief	Notes
1	Welcome	<ul style="list-style-type: none"> <li>Curriculum: <a href="https://medium.com/udacity/term-2-in-depth-on-udacitys-self-driving-car-curriculum-775130aae502#.lra2w8s8m">https://medium.com/udacity/term-2-in-depth-on-udacitys-self-driving-car-curriculum-775130aae502#.lra2w8s8m</a></li> <li>Focus: Sensor fusion, localisation, control</li> </ul>
<b>Sensor Fusion</b>		
2	Introduction to Sensors	<ul style="list-style-type: none"> <li>Types of sensors on a car: stereo camera, traffic signal detection camera, radars, lidars</li> <li>Comparison between Lidar and Radar</li> </ul>
3	Kalman Filters	<ul style="list-style-type: none"> <li>Multivariate Gaussians</li> <li>Intuitive understanding of how KF works - interaction between observable and unobservable states.</li> <li>Linear filter equations</li> </ul>
4	C++ Checkpoint	<ul style="list-style-type: none"> <li>Series of simple exercises. No biggie!</li> </ul>

5	Lidar and Radar Fusion with KF in C++	<ul style="list-style-type: none"> <li>Constant velocity motion model for pedestrian tracking</li> <li>Lidar sensor measurement model</li> <li>Radar sensor measurement model</li> <li>Derivation of Jacobian for radar measurement model</li> <li>Extended Kalman Filter equations</li> <li>Estimating filter performance using RMSE</li> </ul>
6	EKF Project (Project Due 1 May)	<ul style="list-style-type: none"> <li>EKF-based sensor fusion to track objects using lidar and radar</li> <li><a href="https://github.com/cvillas/CarND/tree/master/P6-EKF">https://github.com/cvillas/CarND/tree/master/P6-EKF</a></li> </ul>
7	UKF	<ul style="list-style-type: none"> <li>Constant Turn Rate and Velocity Magnitude (CTRV) process model</li> <li>Generating the sigma-points, applying the process model and computing mean/covariance of the predicted sigma-points.</li> <li>Consistency check using normalised innovation squared</li> <li>Process noise model tuning based on consistency check</li> </ul>
8	UKF Project (Due 15 May)	<ul style="list-style-type: none"> <li>UKF-based sensor fusion for bicycle tracking</li> <li><a href="https://github.com/cvillas/CarND/tree/master/P7-UKF">https://github.com/cvillas/CarND/tree/master/P7-UKF</a></li> </ul>
<b>Localisation</b>		
9	Introduction to Localisation	
10	Localisation Overview	<ul style="list-style-type: none"> <li>Bayes Rule</li> <li>Theorem of total probability</li> </ul>
11	Markov Localisation	<ul style="list-style-type: none"> <li>Probabilistic representations <ul style="list-style-type: none"> <li>Localisation (known map) vs SLAM</li> <li>Likelihood - Observation Model</li> <li>Prior - Motion Model</li> </ul> </li> <li>Markov assumptions and simplification of above probabilistic models</li> <li>Derivation of recursive Bayes' filter for localisation</li> </ul>
12	Motion Models	<ul style="list-style-type: none"> <li>The bicycle model for a car</li> </ul>
13	Particle Filters (PF)	<ul style="list-style-type: none"> <li>Particle filter vs Kalman and histogram filters</li> <li>Conceptual description of steps and their implementation in Python</li> <li>Algorithm for resampling with replacement</li> </ul>
14	Implementation of a PF	<ul style="list-style-type: none"> <li>Implementation of various steps in pseudo-code: initialisation, prediction, data association, weights update, error calculation.</li> <li>Introduction to kidnapped vehicle code</li> </ul>
15	Kidnapped Vehicle Project (Due 5 Jun)	<ul style="list-style-type: none"> <li>Particle filter based localisation</li> <li><a href="https://github.com/cvillas/CarND/tree/master/P8-PF">https://github.com/cvillas/CarND/tree/master/P8-PF</a></li> </ul>
<b>Control</b>		
16	PID Control	<ul style="list-style-type: none"> <li>The individual effects of P, I, and D terms</li> <li>Twiddle algorithm for automatic control gain tuning</li> </ul>
17	PID Controller (Project Due 19 Jun)	<ul style="list-style-type: none"> <li>PID control of steering angle in a simulated car</li> <li><a href="https://github.com/cvillas/CarND/tree/master/P9-PID">https://github.com/cvillas/CarND/tree/master/P9-PID</a></li> <li>Project video: <a href="http://youtube.com/watch?v=xTGGPnw2vfQ">http://youtube.com/watch?v=xTGGPnw2vfQ</a></li> </ul>
18	Vehicle Models	<ul style="list-style-type: none"> <li>Difference between kinematic and dynamic models</li> <li>Review kinematic model</li> <li>Polynomial fitting for smooth desired vehicle trajectories</li> <li>Cross track and orientation error computations</li> <li>Brief intro to dynamic models. No details discussed.</li> </ul>
19	Model Predictive Control	<ul style="list-style-type: none"> <li>Development of a cost function</li> <li>Concept of prediction horizon</li> <li>MPC algorithm development and optimisation (Ipopt and cppad)</li> <li>Incorporating actuation latency into the control algorithm</li> </ul>

20	Model Predictive Control Project (Project Due 3 Jul)	<ul style="list-style-type: none"> <li>Model predictive control of a simulated vehicle around a track, with actuation latency</li> <li><a href="https://github.com/cvilas/CarND/tree/master/P10-MPC">https://github.com/cvilas/CarND/tree/master/P10-MPC</a></li> <li>Project video: <a href="http://youtube.com/watch?v=gweVAQsuLU0">http://youtube.com/watch?v=gweVAQsuLU0</a></li> </ul>
----	--	---

## Term 3

(July 28, 2017 - Nov 6, 2017)

Lesson	Brief	Notes
1	Welcome	Curriculum: <a href="https://medium.com/udacity/term-3-in-depth-on-udacitys-self-driving-car-curriculum-15d03e45d7ea">https://medium.com/udacity/term-3-in-depth-on-udacitys-self-driving-car-curriculum-15d03e45d7ea</a>
2	Search	<ul style="list-style-type: none"> <li>Basic search</li> <li>A* search on grid map</li> <li>Dynamic programming</li> </ul>
3	Prediction	<ul style="list-style-type: none"> <li>Multi-modal nature of driver behaviour prediction</li> <li>Model-based vs data-driven approaches</li> <li>Trajectory clustering and prototype trajectories</li> <li>Frenet coordinates</li> <li>Different types of process models for prediction and <a href="#">their comparison</a>.</li> <li>Hybrid approach to prediction (process model + ML classifier)</li> <li>Naive Bayes classifier example</li> </ul>
4	Behaviour Planning	<ul style="list-style-type: none"> <li>Finite State Machines as an approach to implement behaviour planner <ul style="list-style-type: none"> <li>Strengths and weaknesses of FSM approach</li> <li>Example: highway driving</li> <li>State transition functions</li> </ul> </li> <li>Cost functions and difficulties in designing them properly</li> <li>Scheduling compute time for various modules (behaviour planning, prediction, trajectory planning, localisation, sensor fusion)</li> </ul>
5	Trajectory Generation	<ul style="list-style-type: none"> <li>Definition of motion planning (MP) problem</li> <li>MP algorithm properties - completeness and optimality</li> <li>Classes of MP algorithms - combinatorial, potential fields, optimal control, sampling-based</li> <li>Focus on sampling-based methods - difference between discrete and probabilistic methods</li> <li>Hybrid A* implementation</li> <li>Paper review: <a href="#">Junior: The Stanford entry in the Urban Challenge</a></li> <li>Unstructured vs structured environments</li> <li>Jerk minimising polynomial trajectory generation in Frenet coordinates</li> <li>Paper review: <a href="#">Optimal Trajectory Generation for Dynamic Street Scenarios in a Frenet Frame</a></li> </ul>
6	Project: Path Planning (Due Nov 27)	<ul style="list-style-type: none"> <li>Project integrating trajectory generation with vehicle models, prediction, behaviour planning and control</li> <li><a href="https://github.com/cvilas/CarND/tree/master/P11-PathPlanning">https://github.com/cvilas/CarND/tree/master/P11-PathPlanning</a></li> <li>Project video: <a href="https://www.youtube.com/watch?v=0TvdNeugxE">https://www.youtube.com/watch?v=0TvdNeugxE</a></li> </ul>
7	Intro to electives	
8	<b>Elective: Advanced Deep Learning</b>	
9	Fully Convolutional Networks	
10	Scene Understanding	
11	Inference Performance	
12	Project: Semantic Segmentation (Due Dec 25)	
13	<b>Elective: Functional Safety</b>	

14	Introduction to Functional Safety	<ul style="list-style-type: none"> <li>• Intro to IEC61508 and ISO26262</li> <li>• ISO26262 V-model</li> <li>• Functional safety documentation</li> </ul>
15	Functional Safety: Safety Plan	<ul style="list-style-type: none"> <li>• Safety culture</li> <li>• Safety lifecycle</li> <li>• Safety management roles and responsibilities</li> <li>• Development interface agreement</li> <li>• Confirmation measures</li> </ul>
16	Functional Safety: Hazard Analysis and Risk Assessment	<ul style="list-style-type: none"> <li>• 'Item' definition</li> <li>• Situational analysis</li> <li>• Hazard identification</li> <li>• Hazardous event classification - severity, exposure and controllability</li> <li>• ASIL levels</li> <li>• Definition of safety goals</li> </ul>
17	Functional Safety: Functional Safety Concept	<ul style="list-style-type: none"> <li>• FS requirement attributes - ASIL level, fault tolerant time interval, warning and degradation concept, safe state, verification and validation</li> <li>• ASIL inheritance and decomposition, concept of co-existence</li> </ul>
18	Functional Safety: Technical Safety Concept	
19	Functional Safety at the Software and Hardware Levels	<ul style="list-style-type: none"> <li>• Categories of hardware faults</li> <li>• Safety element out of context (SEooC): <a href="https://www.sae.org/publications/technical-papers/content/2012-01-0033/">https://www.sae.org/publications/technical-papers/content/2012-01-0033/</a></li> <li>• Guidelines for safety in software</li> </ul>
20	Elective Project: Functional Safety (Due: Dec 25)	
21	<b>Autonomous Vehicle Architecture</b>	
22	Introduction to ROS	<ul style="list-style-type: none"> <li>• Nodes, topics, services</li> <li>• rosnodetopics, rostopic list, rostopic info, rostopic echo</li> </ul>
23	Packages and Catkin Workspaces	<ul style="list-style-type: none"> <li>• catkin_init_workspace, catkin_make, roslaunch, rosdep, catkin_create_package</li> </ul>
24	Writing ROS Nodes	<ul style="list-style-type: none"> <li>• Writing new publishers, subscribers, services.</li> <li>• Editing launch files</li> <li>• Logging</li> </ul>
25	Project: Systems Integration (Due Jan 22,'18)	

## References

- Vilas' GitHub repo: <https://github.com/cvilas/CarND>
- A Survey of Motion Planning and Control Techniques for Self-driving Urban Vehicles, <https://arxiv.org/pdf/1604.07446v1.pdf>
- Visualizing and Understanding Convolutional Networks: <https://www.youtube.com/watch?v=ghEmQSxT6tw>
- How the future of autonomous cars will unfold. 16 questions: <https://vimeo.com/198256576>
- Dropout: A Simple Way to Prevent Neural Networks from Overfitting: JMLRdropout.pdf
- Gradient-based learning applied to document recognition (Lenet): [lecun-98.pdf](#)
- Traffic sign recognition with multi-scale convolutional networks: [sermanet-ijcnn-11.pdf](#)
- ImageNet classification with deep convolutional neural networks (AlexNet): [alexnet-12.pdf](#)
- Going deeper with convolutions (GoogLeNet): [GoogLeNet-15.pdf](#)
- Deep residual learning for image recognition (Microsoft Resnet): [resnet-15.pdf](#)
- Very deep convolutional networks for large-scale image recognition (vggnet): [vggnet-15.pdf](#)
- End to end learning for self-driving cars (Nvidia). Original paper: [nvidia\\_end2end-16.pdf](#). Paper explaining how it works: [nvidia\\_end2end2.pdf](#)
- A summary of key recent deep learning architectures: <https://adeshpande3.github.io/adeshpande3.github.io/The-9-Deep-Learning-Papers-You-Need-To-Know-About.html>

- SqueezeNet: AlexNet level accuracy with 50x fewer parameters and < 0.5 MB model size: [squeezeNet-17.pdf](#)
- Batch Normalization: Accelerating deep network training by reducing internal covariate shift: [batch\\_normalisation.pdf](#)
- Histograms of Oriented Gradients for Human Detection: [dalal-cvpr05.pdf](#)
- Nvidia Drive PX2 and what it does: <https://www.youtube.com/watch?v=URmxzxYlmtg>
- Chris Urmson (Google), How a driverless car sees the road: [https://www.ted.com/talks/chris\\_urmson\\_how\\_a\\_driverless\\_car\\_sees\\_the\\_road](https://www.ted.com/talks/chris_urmson_how_a_driverless_car_sees_the_road)
- A comparative study of multiple-model algorithms for maneuvering target tracking: [a-comparative-study-of-multiple-model-algorithms-for-maneuvering-target-tracking.pdf](#)
- Junior: The Stanford entry in the Urban Challenge: [junior\\_stanford-2008.pdf](#)
- Optimal Trajectory Generation for Dynamic Street Scenarios in a Frenet Frame: [werling-optimal-trajectory-generation-10.pdf](#)
- Safety element out of context - A practical approach: <https://www.sae.org/publications/technical-papers/content/2012-01-0033/>
- Pure Pursuit Path Tracking Algorithm: [coulter\\_r\\_craig\\_1992\\_1.pdf](#)
- Localisation in Apollo self-driving car software stack: [apollo\\_localisation-17.pdf](#)