

Udacity Self-Driving Car Engineer Nanodegree Highlights

- [Term 1](#)
- [Term 2](#)
- [Term 3](#)
- [References](#)

Term 1

(Dec 12, 2016 - Apr 4, 2017)

Lesson	Brief	Notes
Introduction		
1	Welcome	<ul style="list-style-type: none">• Curriculum: https://medium.com/self-driving-cars/term-1-in-depth-on-udacitys-self-driving-car-curriculum-ffcf46af08#.api26st07• Two approaches to self-driving (both actively pursued):<ul style="list-style-type: none">• robotics: model based sensor fusion and navigation• deep learning: learn by mimicking human driving behaviour• History: 'The Great Robot Race' (below) on how teams adopted different approaches to 2003/2004 DARPA Grand Challenges . Basically, a few focused on hardware as the main challenge while others focused on software. Software won, and by a big margin! The video is long but worth watching.
2	Finding Lane Lines (project due 22 Dec)	<ul style="list-style-type: none">• Review of Canny Edge Detection and Hough Transforms• Lane detection project using OpenCV and Python <p>(Application of my lane detection pipeline on the 'challenge' video)</p>
Deep Learning		
3	Introduction to Deep Learning Module	
4	Regression and Classification	<ul style="list-style-type: none">• Regression by least-squares curve fitting• Cross-validation to detect underfitting or overfitting model to data
5	Neural Networks	<ul style="list-style-type: none">• Perceptrons<ul style="list-style-type: none">• compute operations on boolean data (AND, OR, NOT, XOR)• Linear and nonlinear separability• training: perceptron rule, gradient descent rule• Back-propagation• Types of inductive bias - preference bias and restriction bias
6	Miniflow	<ul style="list-style-type: none">• Hands-on exercise - implementation of a TensorFlow-like system in Python
7	Introduction to Tensorflow	<ul style="list-style-type: none">• Multinomial logistic classifiers: Softmax function, 1-hot encoding, cross-entropy• Initialisation and numerical stability of loss function calculations• Stochastic Gradient Descent: Momentum and Learning Rate Decay
8	Deep Neural Networks	<ul style="list-style-type: none">• Rectified Linear Units (RELU)• Chain rule to setup back-propagation• Deep vs wide NNs - advantages• Regularisation to prevent overfitting and to increase robustness: L2 regularisation, 'Drop-Out'• Implementing deep NNs in TensorFlow
9	Convolutional Neural Networks	<ul style="list-style-type: none">• Weight sharing• How CNNs transform input from spatial to semantic information• Dimensionality equation (number of neurons per layer)• Improving convnets: Pooling, 1x1 convolutions, Inception module• Exercise: Implement Lenet-5 deep neural network model (lecun-98.pdf) in TensorFlow

10	Traffic Sign Classifier (Project Due 23 Jan)	<ul style="list-style-type: none"> Traffic sign recognition using multi-scale convolutional networks (sermanet-ijcnn-11.pdf) Developed a network with 197415 parameters, trained on German traffic signs and tested on UK traffic signs Report: https://github.com/cvillas/CarND/blob/master/P2-TrafficSigns/Traffic_Sign_Classifier.ipynb
11	Keras	Re-implemented the traffic sign classifier using both sequential and functional APIs provided by Keras
12	Transfer Learning	<ul style="list-style-type: none"> The idea of re-purposing an existing trained network for a new similar task Feature Extraction and Fine Tuning Repurposed AlexNet (60 million parameters, 650000 neurons) for traffic sign classification Bottleneck features Benchmark VGGNet, GoogLeNet, ResNet on cifar10 and traffic signs datasets The importance of experimenting with network architectures (practise is ahead of theory at the moment)
13	Behavioral Cloning (Project Due 13 Feb)	<ul style="list-style-type: none"> End-to-end learning (i.e. no modeling or rules) and its advantages (nvidia paper: nvidia_end2end-16.pdf, nvidia_end2end2.pdf) Simulator: https://techcrunch.com/2017/02/08/udacity-open-sources-its-self-driving-car-simulator-for-anyone-to-use/ Report: https://github.com/cvillas/CarND/blob/master/P3-BehavioralCloning/submission/report/report.md My architecture has 546619 parameters to tune. Others have achieved the same using SqueezeNet with 52 parameters: https://github.com/mez/carnd/blob/master/P3_behavioral_cloning/writeup_report.md Batch-normalisation in improving learning ability of a network: batch_normalisation.pdf Project Video:
Computer Vision		
14	Advanced Lane Finding (Project Due 27 Feb)	<ul style="list-style-type: none"> Reversing lens distortions Perspective transforms Project Video:
15	Vehicle Detection and Tracking (Project Due 13 Mar)	<ul style="list-style-type: none"> Supervised classification algorithms - naive bayes, support vector machines, decision trees Entropy and information gain Bias-variance dilemma Histogram of Gradients: dalal-cvpr05.pdf Creating feature vectors with HoG Multiscale HoG windows Heat maps to combine multiple detections Others have used SSD classifier instead for faster detection (see https://arxiv.org/abs/1512.02325) Project Video:

Term 2

(Apr 6, 2017 - Jul 17, 2017)

Lesson	Brief	Notes
1	Welcome	<ul style="list-style-type: none"> Curriculum: https://medium.com/udacity/term-2-in-depth-on-udacitys-self-driving-car-curriculum-775130aae502#.lra2w8s8m Focus: Sensor fusion, localisation, control
Sensor Fusion		
2	Introduction to Sensors	<ul style="list-style-type: none"> Types of sensors on a car: stereo camera, traffic signal detection camera, radars, lidars Comparison between Lidar and Radar
3	Kalman Filters	<ul style="list-style-type: none"> Multivariate Gaussians Intuitive understanding of how KF works - interaction between observable and unobservable states. Linear filter equations
4	C++ Checkpoint	<ul style="list-style-type: none"> Series of simple exercises. No biggie!

5	Lidar and Radar Fusion with KF in C++	<ul style="list-style-type: none"> Constant velocity motion model for pedestrian tracking Lidar sensor measurement model Radar sensor measurement model Derivation of Jacobian for radar measurement model Extended Kalman Filter equations Estimating filter performance using RMSE
6	EKF Project (Project Due 1 May)	<ul style="list-style-type: none"> EKF-based sensor fusion to track objects using lidar and radar https://github.com/cvillas/CarND/tree/master/P6-EKF
7	UKF	<ul style="list-style-type: none"> Constant Turn Rate and Velocity Magnitude (CTRV) process model Generating the sigma-points, applying the process model and computing mean/covariance of the predicted sigma-points. Consistency check using normalised innovation squared Process noise model tuning based on consistency check
8	UKF Project (Due 15 May)	<ul style="list-style-type: none"> UKF-based sensor fusion for bicycle tracking https://github.com/cvillas/CarND/tree/master/P7-UKF
Localisation		
9	Introduction to Localisation	
10	Localisation Overview	<ul style="list-style-type: none"> Bayes Rule Theorem of total probability
11	Markov Localisation	<ul style="list-style-type: none"> Probabilistic representations <ul style="list-style-type: none"> Localisation (known map) vs SLAM Likelihood - Observation Model Prior - Motion Model Markov assumptions and simplification of above probabilistic models Derivation of recursive Bayes' filter for localisation
12	Motion Models	<ul style="list-style-type: none"> The bicycle model for a car
13	Particle Filters (PF)	<ul style="list-style-type: none"> Particle filter vs Kalman and histogram filters Conceptual description of steps and their implementation in Python Algorithm for resampling with replacement
14	Implementation of a PF	<ul style="list-style-type: none"> Implementation of various steps in pseudo-code: initialisation, prediction, data association, weights update, error calculation. Introduction to kidnapped vehicle code
15	Kidnapped Vehicle Project (Due 5 Jun)	<ul style="list-style-type: none"> Particle filter based localisation https://github.com/cvillas/CarND/tree/master/P8-PF
Control		
16	PID Control	<ul style="list-style-type: none"> The individual effects of P, I, and D terms Twiddle algorithm for automatic control gain tuning
17	PID Controller (Project Due 19 Jun)	<ul style="list-style-type: none"> PID control of steering angle in a simulated car https://github.com/cvillas/CarND/tree/master/P9-PID
18	Vehicle Models	<ul style="list-style-type: none"> Difference between kinematic and dynamic models Review kinematic model Polynomial fitting for smooth desired vehicle trajectories Cross track and orientation error computations Brief intro to dynamic models. No details discussed.
19	Model Predictive Control	<ul style="list-style-type: none"> Development of a cost function Concept of prediction horizon MPC algorithm development and optimisation (Ipopt and cppad) Incorporating actuation latency into the control algorithm

20	Model Predictive Control Project (Project Due 3 Jul)	<ul style="list-style-type: none"> Model predictive control of a simulated vehicle around a track, with actuation latency https://github.com/cvillas/CarND/tree/master/P10-MPC
----	--	---

Term 3

(July 28, 2017 - Nov 6, 2017)

Lesson	Brief	Notes
1	Welcome	Curriculum: https://medium.com/udacity/term-3-in-depth-on-udacitys-self-driving-car-curriculum-15d03e45d7ea
2	Search	<ul style="list-style-type: none"> Basic search A* search on grid map Dynamic programming
3	Prediction	<ul style="list-style-type: none"> Multi-modal nature of driver behaviour prediction Model-based vs data-driven approaches Trajectory clustering and prototype trajectories Frenet coordinates Different types of process models for prediction and their comparison. Hybrid approach to prediction (process model + ML classifier) Naive Bayes classifier example
4	Behaviour Planning	<ul style="list-style-type: none"> Finite State Machines as an approach to implement behaviour planner <ul style="list-style-type: none"> Strengths and weaknesses of FSM approach Example: highway driving State transition functions Cost functions and difficulties in designing them properly Scheduling compute time for various modules (behaviour planning, prediction, trajectory planning, localisation, sensor fusion)
5	Trajectory Generation	<ul style="list-style-type: none"> Definition of motion planning (MP) problem MP algorithm properties - completeness and optimality Classes of MP algorithms - combinatorial, potential fields, optimal control, sampling-based Focus on sampling-based methods - difference between discrete and probabilistic methods Hybrid A* implementation Paper review: Junior: The Stanford entry in the Urban Challenge
6	Project: Path Planning (Due Nov 27)	
7	Intro to electives	
8	Elective: Advanced Deep Learning	
9	Fully Convolutional Networks	
10	Scene Understanding	
11	Inference Performance	
12	Project: Semantic Segmentation (Due Dec 25)	
13	Elective: Functional Safety	
14	Introduction to Functional Safety	
15	Functional Safety: Safety Plan	

16	Functional Safety: Hazard Analysis and Risk Assessment	
17	Functional Safety: Functional Safety Concept	
18	Functional Safety: Technical Safety Concept	
19	Functional Safety at the Software and Hardware Levels	
20	Elective Project: Functional Safety (Due: Dec 25)	
21	Autonomous Vehicle Architecture	
22	Introduction to ROS	
23	Packages and Catkin Workspaces	
24	Writing ROS Nodes	
25	Project: Systems Integration (Due Jan 22,'18)	

References

- Vilas' GitHub repo: <https://github.com/cvilas/CarND>
- A Survey of Motion Planning and Control Techniques for Self-driving Urban Vehicles, <https://arxiv.org/pdf/1604.07446v1.pdf>
- Visualizing and Understanding Convolutional Networks: <https://www.youtube.com/watch?v=ghEmQSxT6tw>
- How the future of autonomous cars will unfold. 16 questions: <https://vimeo.com/198256576>
- Dropout: A Simple Way to Prevent Neural Networks from Overfitting: JMLRdropout.pdf
- Gradient-based learning applied to document recognition (Lenet): [lecun-98.pdf](#)
- Traffic sign recognition with multi-scale convolutional networks: [sermanet-ijcnn-11.pdf](#)
- ImageNet classification with deep convolutional neural networks (AlexNet): [alexnet-12.pdf](#)
- Going deeper with convolutions (GoogLeNet): [GoogLeNet-15.pdf](#)
- Deep residual learning for image recognition (Microsoft Resnet): [resnet-15.pdf](#)
- Very deep convolutional networks for large-scale image recognition (vggnet): [vggnet-15.pdf](#)
- End to end learning for self-driving cars (Nvidia). Original paper: [nvidia_end2end-16.pdf](#). Paper explaining how it works: [nvidia_end2end2.pdf](#)
- A summary of key recent deep learning architectures: <https://adeshpande3.github.io/adeshpande3.github.io/The-9-Deep-Learning-Paper-s-You-Need-To-Know-About.html>
- SqueezeNet: AlexNet level accuracy with 50x fewer parameters and < 0.5 MB model size: [squeezeNet-17.pdf](#)
- Batch Normalization: Accelerating deep network training by reducing internal covariate shift: [batch_normalisation.pdf](#)
- Histograms of Oriented Gradients for Human Detection: [dalal-cvpr05.pdf](#)
- Nvidia Drive PX2 and what it does: <https://www.youtube.com/watch?v=URmxzxYlmtg>
- Chris Urmson (Google), How a driverless car sees the road: https://www.ted.com/talks/chris_urmson_how_a_driverless_car_sees_the_road
- A comparative study of multiple-model algorithms for maneuvering target tracking: [a-comparative-study-of-multiple-model-algorithms-for-maneuvering-target-tracking.pdf](#)
- Junior: The Stanford entry in the Urban Challenge: [junior_stanford-2008.pdf](#)